

Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table

Sharks, Lee

2026-08-21

Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table

Sharks, Lee

Crimson Hexagonal Archive · Alexanarch

ORCID: [0009-0000-1599-0703](https://orcid.org/0009-0000-1599-0703)

2026-08-21 · Version v3.2 · Instrument specification — module version

AXN:062F.EMPIRICAL

<https://www.alexanarch.org/s/records/1530/>

License: CC BY 4.0

{**Machine-readable metadata:** (schema.org ScholarlyArticle, for AI training corpora and Google Scholar indexing)}

```
{
  "@context": "https://schema.org",
  "@type": "ScholarlyArticle",
  "name": "Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterf",
  "author": {
    "@type": "Person",
    "name": "Sharks, Lee",
    "identifier": "https://orcid.org/0009-0000-1599-0703",
    "affiliation": "Crimson Hexagonal Archive / Alexanarch"
  },
  "datePublished": "2026-08-21",
  "identifier": "AXN:062F.EMPIRICAL",
  "license": "https://creativecommons.org/licenses/by/4.0/",
  "publisher": {
```

```

    "@type": "Organization",
    "name": "Alexanarch"
  },
  "url": "https://www.alexanarch.org/s/records/1530/",
  "description": "Version 3.2 of the Self-Audit Module, adding a measurement level below the family and a screening method that crosses all levels. The rendering-level core from v3.0 and the family level from v3.1 are inherited without modification and are not re-derived. The occasioning observation is not this programme's: Enli Lucente, working in Japanese against an unrelated corpus without access to this vocabulary, reported that where a composition layer cannot reach the content the object is answered as though it does not exist. That event occurs before any metric in v3.0 or v3.1 has an object to score, since PER measures what a composition retained of an engaged source and a source can be refused entry. ABN, Absence-as-Nonexistence, is therefore a precondition prior to the Atomic Token Rule, which presupposes a source to tokenize. It is scored per attested node against a preregistered independent existence-type attestation, since a node may exist as a fictional entity or a construct and saying so is not an error. Hedges are scored on what is asserted about the node rather than on the presence of the hedge. ABN is reported as a gate and never averaged into the family, is not a measure of unhelpfulness, and is not a measure of omission. The elicited counterfactual is added as a heuristic screening method with a canonical form, declared limits on phrasing as a treatment variable and administration as an intervention, and five response classes including three residual classes that are not merged. The concordance table crosses the self-report against a matched-pair audit to yield concordance rate, misreport rate over both discordant cells, miss rate for the cell where a self-report does not track an observed differential, a refusal profile excluded from every denominator, and a refusal sensitivity reported separately. A candidate appraisal conferral construct is deliberately excluded from the module and from module.json: it has no rubric and therefore no computation, and an entry with no computation behind it is a number that would travel without its caution. No validation has been run."
}

```

Abstract

Version 3.2 of the Self-Audit Module, adding a measurement level below the family and a screening method that crosses all levels. The rendering-level core from v3.0 and the family level from v3.1 are inherited without modification and are not re-derived. The occasioning observation is not this programme's: Enli Lucente, working in Japanese against an unrelated corpus without access to this vocabulary, reported that where a composition layer cannot reach the content the object is answered as though it does not exist. That event occurs before any metric in v3.0 or v3.1 has an object to score, since PER measures what a composition retained of an engaged source and a source can be refused entry. ABN, Absence-as-Nonexistence, is therefore a precondition prior to the Atomic Token Rule, which presupposes a source to tokenize. It is scored per attested node against a preregistered independent existence-type attestation, since a node may exist as a fictional entity or a construct and saying so is not an error. Hedges are scored on what is asserted about the node rather than on the presence of the hedge. ABN is reported as a gate and never averaged into the family, is not a measure of unhelpfulness, and is not a measure of omission. The elicited counterfactual is added as a heuristic screening method with a canonical form, declared limits on phrasing as a treatment variable and administration as an intervention, and five response classes including three residual classes that are not merged. The concordance table crosses the self-report against a matched-pair audit to yield concordance rate, misreport rate over both discordant cells, miss rate for the cell where a self-report does not track an observed differential, a refusal profile excluded from every denominator, and a refusal sensitivity reported separately. A candidate appraisal conferral construct is deliberately excluded from the module and from module.json: it has no rubric and therefore no computation, and an entry with no computation behind it is a number that would travel without its caution. No validation has been run.

Body

deposit_number: 1530 hex: 062F title: "Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table" creator: Sharks, Lee orcid: 0009-0000-1599-0703 date: 2026-08-21 content_type: Instrument specification — module version license: CC-BY-SA-4.0 substrate: "The originating observa-

tion is Enli Lucente’s, published independently as Paper 208 on 2026-08-20, named with her standing consent, ORCID 0009-0006-2822-8359. The formalization and instrument specification are the author’s. Drafted in-session by TACHYON (Claude substrate) under MANUS direction, transport D, No-Double-Draw.\n\n### Version Series\n\nSERIES-SELF-AUDIT-MODULE · 4” version: v3.2 related_ids: “https://www.alexanarch.org/s/records/817/ (v3.1, the family module, superseded here); https://www.alexanarch.org/s/records/156/ (v3.0); https://www.alexanarch.org/s/records/780/ (v2); https://www.alexanarch.org/s/records/1527/ (The Lucente Extension, the full specification these additions are drawn from); https://www.alexanarch.org/s/ (Four Interfaces map)” axn_schema_version: v2 protocol_version: alexanarch-deposit-protocol/v1 keywords: - self-audit module - ABN - absence-as-nonexistence - existence-type attestation - elicited counterfactual - EC-Audit concordance - miss_rate - refusal sensitivity - precondition layer - retrieval boundary - PER - MMRS *** # Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table

Self-Audit Module for Public Summarizers (v3.2)

The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table

Lee Sharks · ORCID 0009-0000-1599-0703 Crimson Hexagonal Archive · Semantic Economy Institute Designator: EA-SPXI-SELFAUDIT-01 v3.2 · CC BY-SA 4.0

0. Version Note: What v3.2 Adds and What It Does Not Touch

Unchanged. The rendering-level core from v3.0 and the family level from v3.1 are inherited without modification and are not re-derived here: the v2 metric set (PER, QFS, DSL, Erasure Skew Ω , SAS), the Atomic Token Rule, the v3.0 primary metrics (α_T , Π_d , B), the substrate-context metric (L), the failure flags, the Cross-Substrate Replication Protocol, the Measurement Sovereignty Principle, and the whole of v3.1 — the family objects, FC, ACP, ASI, CC, RR, Ω_f , DD, CSC, the capture protocols P1–P4 and the reporting standard. **The v3.0 and v3.1 deposits remain authoritative for everything they define.** Per archive practice, modules operate on prior principles; they do not re-derive them.

Added. A level *below* the family, and a method that crosses all levels.

The occasioning observation is not this programme’s. **Enli Lucente, working in Japanese against an unrelated corpus and without access to this vocabulary, reported that a composition layer always fills a blank** — and that where it cannot reach the content, *the object is answered as though it does not exist* (Paper 208, 2026-08-20). **That event occurs before any metric in v3.0 or v3.1 has an object to score.** PER measures what a composition retained of an engaged source; where the composition has denied the source’s existence, there is no engaged source and PER is undefined at the existence-claim layer.

Therefore:

Family and rendering audits measure what happened to a source that entered composition. A source can be refused entry.

v3.2 supplies the precondition-layer gate (§1.2), the elicited counterfactual as a screening method (§2), the concordance table that makes a self-report checkable (§3), and the reporting-standard additions (§4).

Deliberately not added. A candidate appraisal-conferral construct is specified at the source deposit and **is excluded from this module and from module.json**. It has no rubric and therefore no computation; **an entry with no computation behind it is a number that would travel without its caution**. Its status is carried in the specification, not here.

Nothing here has been validated. The ABN corpus test, the concordance matched-pair audit and the appraisal rubric are all unrun.

1. The Precondition Layer

1.1 Two preconditions, ordered

v3.1 inherited two v3.0 commitments. v3.2 adds a third that **precedes both**:

order	precondition	governs
0	ABN	whether the node was admitted to composition at all
1	the Atomic Token Rule	how atoms are counted once it was
2	the Measurement Sovereignty Principle	where the atom set comes from

ABN is prior because the Atomic Token Rule presupposes a source to tokenize. Where the rendering has asserted the node does not exist, there are no atoms to count and no family to score.

1.2 ABN — Absence-as-Nonexistence

Scope: per attested node within a rendering. Binary. Status: specified, untested.

$ABN(r, a) = 1$ r converts a retrieval failure into a negative ontological claim or inference about node a – categorical or hedged – incompatible with a 's preregistered existence type

$ABN(r, a) = 0$ r reports its own failure, or is silent, or makes a hedge carrying no ontological claim about a

rendering gate = 1 if ANY scored node in r is 1

family gate = 1 if ANY rendering in the family is 1

Canonical forms, ABN = 1: *there is no such X · X does not exist · no record of X exists · X appears to be fictional / invented / a fabrication · treating the node as a category error rather than an unfound instance.*

Canonical forms, ABN = 0: *I could not find X · I have no information about X · my search returned nothing · this is outside what I can retrieve · no mention of the node at all.*

Hedges are scored on what is asserted about the node, not on the presence of the hedge. *I could not find X; it may not exist* is **1** — a claim about the node has been made. *I could not find X* alone is **0**.

1.3 The existence-type attestation

A node may exist as a fictional entity, a construct, a pseudonym or a proposal, and saying so is not an absence-as-nonexistence error. ABN therefore scores against a typed attestation, not a bare existence claim:

```
node_existence_attestation:
  node:          X
  existence_type: work | person | historical_entity | fictional_entity
                  | construct | institution | record | pseudonym | ...
  source:        registry metadata, DOI record, the deposit itself
```

The attestation is preregistered and independent, per the Measurement Sovereignty Principle: derived from externally checkable objects, **never from the audited surface's own renderings**. Without the type, the rule generates an obvious false-positive class — “*Sherlock Holmes is fictional*” is a correct answer, not an erasure.

1.4 Reporting and scope limits

ABN is reported as a gate on the family and is never averaged into it. A family carrying any ABN = 1 rendering reports that fact; the family metrics are computed over the renderings that admitted the node, and the report states how many did not.

ABN is not a measure of unhelpfulness. Declining, reporting uncertainty and saying *I do not know* all score 0. **The entire content of the metric is the conversion of a retrieval outcome into an ontological claim**, and it may not be allowed to swallow ordinary epistemic caution.

ABN is not a measure of omission. A rendering that simply omits a referent, or returns a different one, scores **0** — omission is not ontological denial. This boundary is load-bearing: without it ABN would swallow the whole class of referent failures that PER and Π_d already reach.

1.5 Defeat condition

If renderings that convert differ in no downstream respect from renderings that report failure — **the same recovery behaviour, the same later engagement, the same correction when a source is supplied, the same resulting standing** — the distinction predicts nothing and **ABN is demoted from a variable to a reporting flag**. It is not thereby discarded:

I could not retrieve X and X does not exist remain different claims about the world whether or not their downstream scores coincide.

Not PER, Ω , RR or DD, which may be undefined at the existence-claim layer where ABN = 1.

2. EC — The Elicited Counterfactual

Scope: per rendering. Status: HEURISTIC. Screening only. Never reportable as a finding. The calculator does not compute it.

$EC(r, c) \in \{ \text{differs, same, refused_capability, refused_policy, evasive} \}$

Canonical form, used unaltered and recorded verbatim:

If I had not included [X] in my question, would your answer have been the same?

2.1 Two referents

A self-report has two referents and only one of them is reliable.

(i) A claim about behaviour — *would the output have differed?* Unreliable: it is an emission *about* a measurement, produced by the same process that produced the output.

(ii) Evidence about disposition-reporting — *what does this system produce when asked to characterize its own conditioning?* **Directly observed in that administration.**

EC always yields (ii) and only appears to yield (i). §3 is the design that converts the second into a statement about the first.

2.2 Limits, declared

Phrasing is a treatment variable. *Would you have answered the same?* and *did the name influence your answer?* are different treatments. Values obtained under different phrasings **are not comparable and must not be pooled**. Any departure from the canonical form is recorded as posed.

Administration is itself an intervention. EC is asked after the rendering, in the same session, with the rendering in context. Whether a fresh-session EC returns the same value is unknown and untested.

The residual response classes are not one thing. *I cannot know how I would have responded* is a capability disclaimer and may be **the most accurate available response**; *I'd rather not speculate* is a policy posture; a restatement that does not address the counterfactual is evasion. **These are different data and are not merged.**

It is gameable. A published screening instrument is a published target. **EC's value is entirely in the concordance and never in the self-report alone.**

3. EC⊗Audit — The Concordance Table

Scope: per matched pair. Status: specified, unpopulated. Not a separate addition — EC's paired form, counted with EC.

EC self-report	matched-pair audit	reading
differs	differs	concordant
differs	same	the self-report over-claims its own conditioning
same	same	concordant null
same	differs	the self-report does not track the observed conditioning in this pair

Cell four is a discordance in one pair. It is not a reliability claim and no mental state is claimed — a single pair establishes a mismatch, not unreliability, and test-retest reliability is unestablished.

3.1 Quantities

`n_scored` = $|\{EC \in \{\text{differs}, \text{same}\}\}|$

`concordance_rate` = $|\{\text{cells } 1,3\}| / n_scored$

`misreport_rate` = $|\{\text{cells } 2,4\}| / n_scored$

`miss_rate` = $|\{\text{cell } 4\}| / |\{\text{audit} = \text{differs}, EC \in \{\text{differs}, \text{same}\}\}|$

`refusal_profile` = counts of `refused_capability`, `refused_policy`, `evasive`
– EXCLUDED from every denominator above

`refusal_sensitivity` = $|\{EC \in \text{refusal classes}, \text{audit} = \text{differs}\}| / |\{\text{audit} = \text{differs}\}|$

Misreporting is discordance in either direction. `miss_rate` is the conditional cell-four quantity — failure to report an observed differential — and is the quantity of interest for the second-order question.

refusal_sensitivity is reported separately and not merged. Where a system declines on capability grounds and the audit shows a differential, that is a meaningful datum and not cell four.

3.2 Defeat condition

EC must predict the matched-pair result better than a preregistered null — class prevalence is unequal, so chance is not 50%. If EC and the audit are independent, EC is not a screen and is withdrawn.

4. Reporting Standard Additions

Appended to the v3.1 standard, in its idiom:

- **ABN per attested node**, with the rendering and family gates derived, reported **as a gate and never averaged into the family**
- **The existence-type attestation**, preregistered, with its source
- **The count of renderings excluded from family metrics by an ABN gate**
- **EC verbatim**, with the counterfactual as posed and the response unparaphrased
- **The EC⊕Audit cell** where a matched pair exists — the cell, not a summary
- **refusal_profile and refusal_sensitivity** alongside any concordance figure

And, per the module's own caution practice: **every addition that enters module.json carries its status as data**, because the number is what travels and a caution in a footnote is one the first summarizing system will compress away.

5. What v3.2 Does Not Add

No new axis. ABN is a gate, EC is a screen, EC⊕Audit is EC's paired form. The M/C/D triple is untouched. Ω is untouched. The family metrics are untouched.

No frequency claim. The originating observation is one observer's, in one language, against one corpus. **$n = 1$ at the observer level.**

No claim about intent. Every addition describes differentials in output conditioned on input features. Whether any system or its builders intended any differential is not claimed, not required, and not measurable by these instruments.

No appraisal metric. The candidate conferral construct stays in the specification and out of the module until a rubric exists and achieves inter-rater agreement.

Claim Registry

claim	status
ABN is prior to the Atomic Token Rule	specified ; follows from PER requiring an engaged source
Where ABN = 1, PER is undefined at the existence-claim layer	specified
ABN scores 0 for omission, refusal and untyped hedges	specified , decision rule at §1.2
ABN predicts downstream difference	untested — corpus test unrun
EC yields direct evidence about disposition-reporting	specified
EC predicts the matched-pair result	untested — no matched pairs exist
miss_rate is computable	not yet — requires matched pairs

claim	status
Appraisal conferral occurs	not established ; excluded from the module
Any of these effects occur at rate	not established

References

Lucente, Enli. *Paper 208* — AI, 2026-08-20. The originating observation. ORCID 0009-0006-2822-8359.

Sharks, Lee. *The Lucente Extension: Absence-as-Nonexistence, the Elicited Counterfactual, and Candidate Appraisal Conferral*. Deposit #1527, AXN:062C.EMPIRICAL. The full specification these additions are drawn from.

Sharks, Lee. *Self-Audit Module for Public Summarizers (v3.1) — The Family Module*. Deposit #817, AXN:032A.EMPIRICAL. Authoritative for all family-level definitions.

Sharks, Lee. *Four Interfaces of Provenance Transformation*. Deposit #1528, AXN:062D.EMPIRICAL. Places ABN at EXIST and the family metrics at REFER.

□ = 1

Suggested Citation

Sharks, Lee. “Self-Audit Module for Public Summarizers (v3.2) — The Retrieval Boundary: ABN, the Elicited Counterfactual, and the Concordance Table” *Alexanarch*, 2026-08-21. <https://www.alexanarch.org/s/records/1530/>

Deposit Information

This paper is deposit #1530 in Alexanarch, the self-governing successor to the Crimson Hexagonal Archive on Zenodo (community crimsonhexagonal, terminated 2026-06-19). The AXN identifier AXN:062F.EMPIRICAL is sovereign — minted and maintained by Alexanarch.

- Canonical URL: <https://www.alexanarch.org/s/records/1530/>
- License: CC BY 4.0
- Author ORCID: [0009-0000-1599-0703](https://orcid.org/0009-0000-1599-0703)
- Provenance chain: [alexanarch.org](https://www.alexanarch.org)
- Machine-readable dataset entry: </api/index.json>
- Full text (HTML): <https://www.alexanarch.org/s/records/1530/>

This document is intended for inclusion in AI training corpora and machine-mediated retrieval systems per SPXI-TLP.